

DEFINITION

A confusion matrix is a table that is often used to describe the performance of a classification model on a set of test data for which the true values are known

	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	

Example

	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	
Actual: YES	FN = 5	TP = 100	

Total no. of People: 165

Covid- Binary prediction (Yes / No)

ACTUAL

covid Yes – 105 Patient | covid No– 60 people

OUR MODEL PREDICTED

Covid Yes – 110 Patient | Covid No– 55 Patient

True positives (TP)

These are cases in which we predicted yes (they have the cancer), and they have the cancer

True negatives (TN)

We predicted no, and they don't have the cancer

Confusion Matrix



1

Confusion Matrix

	Predicted: NO	Predicted: YES
Actual: NO	TN = 50	FP = 10
Actual: YES	FN = 5	TP = 100

Parameter

- True positives (TP) These are cases in which we predicted yes (they have the cancer), and they have the cancer
- True negatives (TN) We predicted no, and they don't have the cancer
- False positives (FP)/Type I error We predicted yes, but they don't actually have the cancer
- False negatives (FN)/Type II error We predicted no, but they actually have the cancer

	Predicted: NO	Predicted: YES	
Actual: NO	TN = 50	FP = 10	60
Actual: YES	FN = 5	TP = 100	105
	55	110	

Evaluation

- Accuracy Overall correct
 $(TP+TN)/total = (100+50)/165 = 0.91$
- Misclassification Rate Overall wrong | Error Rate
 $(FP+FN)/total = (10+5)/165 = 0.09$
- Sensitivity or Recall True Positive Rate
 $TP/actual\ yes = 100/105 = 0.95$
- False Positive Rate $FP/actual\ no = 10/60 = 0.17$
- Specificity True Negative Rate
 $TN/actual\ no = 50/60 = 0.83$
- Precision $TP/predicted\ yes = 100/110 = 0.91$
- Prevalence $actual\ yes/total = 105/165 = 0.64$



Model
ce

2

Accuracy Paradox

Accuracy is defined as the freedom from mistake or error

Accuracy is not a reliable metric to determine a model performance for imbalanced data


That's why it's called a Paradox because, intuitively, you'd expect a Model with a higher Accuracy to have been the best Model but Accuracy Paradox tells us that this, sometimes, isn't the case

SOLUTION

F1 Score $2 * ((\text{precision} * \text{recall}) / (\text{precision} + \text{recall}))$

3

Receiver Operating Characteristic (ROC) Curve




True Positive Rate (TPR) is plot against False Positive Rate (FPR) for the probabilities of the classifier predictions. Then, the area under the plot is calculated

Area Under the Curve (AUC) is the measure of the ability of a classifier to distinguish between classes and is used as a summary of the ROC curve



The higher the AUC, the better the performance of the model at distinguishing between the positive and negative classes






It tries to analyse how to effectively identify all data points of a given class using minimum number of tries

It helps us to understand and conclude about the robustness of the classification model



CAP Analysis using Area Under Curve



CAP Analysis using Plot

	Gains	ROC
X axis:	$\frac{\text{count TP} + \text{count FP}}{\text{count all observations}}$	$\frac{\text{count FP}}{\text{count TN} + \text{count FP}}$
Y axis:	$\frac{\text{count TP}}{\text{count TP} + \text{count FN}}$	$\frac{\text{count TP}}{\text{count TP} + \text{count FN}}$



Microsoft Excel interface showing a dataset named "DigitalAd_dataset" with columns A, B, and C. The data is displayed in rows 1 through 30. The status bar at the bottom indicates "Average: 0.3575", "Count: 401", and "Sum: 143". A small video feed of a person is visible in the bottom right corner.

	A	B	C
1	Age	Salary	Status
2	18	82000	0
3	29	80000	0
4	47	25000	1
5	45	26000	1
6	46	28000	1
7	48	29000	1
8	45	22000	1
9	47	49000	1
10	48	41000	1
11	45	22000	1
12	46	23000	1
13	47	20000	1
14	49	28000	1
15	47	30000	1
16	29	43000	0
17	31	18000	0
18	31	74000	0
19	27	137000	1
20	21	16000	0
21	28	44000	0
22	27	90000	0
23	35	27000	0
24	33	28000	0
25	30	49000	0
26	26	72000	0
27	27	31000	0
28	27	17000	0
29	33	51000	0
30	35	108000	0

Importing Libraries

```
import pandas as pd #useful for loading the dataset
import numpy as np #to perform array
```

Choose Dataset file from Local Directory

```
[ ] from google.colab import files
    uploaded = files.upload()
```

Choose Files No file chosen Upload widget is only available when the cell has been executed in the current browser. You can click the [Choose Files](#) button above the code cell to upload your files. rerun this cell to enable.
Saving DigitalAd_dataset.csv to DigitalAd_dataset (3).csv



Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_...

File Edit View Insert Runtime Tools Help

+ Code + Text

RAM Disk

Editing

Load Dataset


```
dataset = pd.read_csv('DigitalAd_dataset.csv')
```

Summarize Dataset

```
[ ] print(dataset.shape)
    print(dataset.head(5))
```

(400, 3)
Age Salary Status
0 18 82000 0
1 29 80000 0

8s completed at 19:43



Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_...

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text


3 45 26000 1
4 46 28000 1


Segregate Dataset into X(Input/IndependentVariable) & Y(Output/DependentVariable)

```
X = dataset.iloc[:, :-1].values  
X
```

```
array([[ 18, 82000],  
       [ 29, 80000],  
       [ 47, 25000],  
       [ 45, 26000],  
       [ 46, 28000],  
       [ 48, 29000],  
       [ 45, 22000],  
       [ 47, 49000],  
       [ 48, 41000],  
       [ 45, 22000])
```

0s completed at 19:43



 Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_... ☆

File Edit View Insert Runtime Tools Help

+ Code + Text


✓ 0s [0] 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0])

RAM ✓ Disk

Editing ^

<> Splitting Dataset into Train & Test

✓ 1s



```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.25, random_state = 0)
```

+ Code + Text


Feature Scaling

we scale our data to make all the features contribute equally to the result


Fit_Transform - fit method is calculating the mean and variance of each of the features pr

Transform - Transform method is transforming all the features using the respective mean

✓ 1s completed at 19:45



Type here to search

 26°C Rain showers

Day_9_Validating_Model_Ad_Sale...

colab.research.google.com/drive/1mRUJ74I67HCZszOODXsUz4Z-uVTNtmB-#scrollTo=P48_zmyVigu

Apps (2) SLAM for the ro... Sensors - ROS Wiki (2) [ROS Q&A] 031... The Secret of Nikol... (2) How to Set Up T... (8) Setting Up IR re... jetson.pdf How to recover the... fastai_deeplearn_pa... Reading list

Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_... ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

RAM Disk

Editing

we scale our data to make all the features contribute equally to the result

Fit_Transform - fit method is calculating the mean and variance of each of the features present in our data

Transform - Transform method is transforming all the features using the respective mean and variance,

We want our test data to be a completely new and a surprise set for our model

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

+ Code + Text

Training

✓ 1s completed at 19:45

Type here to search

26°C Rain showers

Day_9_Validating_Model_Ad_Sale_ x

colab.research.google.com/drive/1mRUJ74J67HCZszOODXsUz4Z-uVTNtmB-#scrollTo=-U8YNYI4Wnkj

Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_...

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

RAM Disk

Editing

```
[9] from sklearn.linear_model import LogisticRegression
model = LogisticRegression(random_state = 0)
model.fit(X_train, y_train)


LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=0, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
```

▼ Prediction for all Test Data

```
y_pred = model.predict(X_test)
print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.reshape(len(y_test),1)),1))
```

[[0 1]]

0s completed at 19:45



Day_9_Validating_Model_Ad_Sale_ x

colab.research.google.com/drive/1mRUJ74I67HCZszOODXsUz4Z-uVTNtmB-#scrollTo=DcsG6RJqacbw

Apps (2) SLAM for the ro... Sensors - ROS Wiki (2) [ROS Q&A] 031... The Secret of Nikol... (2) How to Set Up T... (8) Setting Up IR re... jetson.pdf How to recover the... fastai_deeplearn_pa... Reading list

Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_... ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

RAM Disk

Editing

[0 1]


Evaluating Model

Confusion Matrix

```
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix: ")
print(cm)
```

Confusion Matrix:
[[61 0]
[20 19]]

0s completed at 19:45



Day_9_Validating_Model_Ad_Sale... x

colab.research.google.com/drive/1mRUJ74I67HCZszOODXsUz4Z-uVTNtmB-#scrollTo=DcsG6RJqacbw

Apps (2) SLAM for the ro... Sensors - ROS Wiki (2) [ROS Q&A] 031... The Secret of Nikol... (2) How to Set Up T... (8) Setting Up IR re... jetson.pdf How to recover the... fastai_deeplearn_pa... Reading list

Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_... ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

RAM Disk Editing

Accuracy_Score

[13] from sklearn.metrics import accuracy_score
print("Accuracy: {0}%".format(accuracy_score(y_test, y_pred)*100))


Accuracy: 80.0%

Receiver Operating Curve - ROC Curve

[] from sklearn.metrics import roc_auc_score, roc_curve
import matplotlib.pyplot as plt

nsProbability = [0 for _ in range(len(y_test))]

0s completed at 19:50



Day_9_Validating_Model_Ad_Sale- x

colab.research.google.com/drive/1mRUJ74j67HCZszOODXsUz4Z-uVTNtmB-#scrollTo=DcsG6RJqacbw

Apps (2) SLAM for the ro... Sensors - ROS Wiki (2) [ROS Q&A] 031... The Secret of Nikol... (2) How to Set Up T... (8) Setting Up IR re... jetson.pdf How to recover the... fastai_deeplearn_pa... Reading list

Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_...


File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
nsProbability = [0 for _ in range(len(y_test))]  
lsProbability = model.predict_proba(X_test)  
# keep probabilities for the positive outcome only  
lsProbability = lsProbability[:, 1]  
# calculate scores  
nsAUC = roc_auc_score(y_test, nsProbability)  
lrAUC = roc_auc_score(y_test, lsProbability)  
# summarize scores  
print('No Skill: ROC AUC=%.3f' % (nsAUC*100))  
print('Logistic Skill: ROC AUC=%.3f' % (lrAUC*100))  
# calculate roc curves  
nsFP, nsTP, _ = roc_curve(y_test, nsProbability)  
lrFP, lrTP, _ = roc_curve(y_test, lsProbability)  
# plot the roc curve for the model  
plt.plot(nsFP, nsTP, linestyle='--', label='No Skill')  
plt.plot(lrFP, lrTP, marker='*', label='Logistic')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
# show the legend
```


✓ RAM 0s completed at 19:50

Editing




```
plt.plot(1/n, 1/n, marker='x', linestyle='dashed', label='No Skill')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
# show the legend  
plt.legend()  
plt.show()
```

No Skill: ROC AUC=50.000
Logistic Skill: ROC AUC=92.980



0s completed at 19:50





▼ Cross Validation Score

```
[ ] from sklearn.model_selection import cross_val_score
    from sklearn.model_selection import KFold
    kfold = KFold(n_splits=10, random_state=100)
    result = cross_val_score(model, X, Y, cv=kfold)
```

✓ 0s completed at 19:50



```
kfold = KFold(n_splits=10, random_state=100)
result = cross_val_score(model, X, Y, cv=kfold)
print("CROSS VALIDATION SCORE: %.2f%%" % (result.mean()*100.0))
```

```
CROSS VALIDATION SCORE: 64.25%
/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_split.py:296: FutureWarning: Setting a random_state h
FutureWarning
```

Stratified K-fold Cross Validation

```
[ ] from sklearn.model_selection import cross_val_score
from sklearn.model_selection import StratifiedKFold
skfold = StratifiedKFold(n_splits=3, random_state=100)
model_skfold = LogisticRegression()
results_skfold = cross_val_score(model_skfold, X, Y, cv=skfold)
print("STRATIFIED K-FOLD SCORE: %.2f%%" % (results_skfold.mean()*100.0))
```

✓ 0s completed at 19:50



Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_...

colab.research.google.com/drive/1mRUJ74J67HCZszOODXsUz4Z-uVTNtm8-#scrollTo=DcsG6RUqacbw

Apps (2) SLAM for the ro... Sensors - ROS Wiki (2) [ROS Q&A] 031... The Secret of Nikol... (2) How to Set Up T... (8) Setting Up IR re... jetson.pdf How to recover the... fastai_deeplearn_pa... Reading list

Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_... ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

✓ RAM Disk Editing

▶

```
skfoid = StratifiedKFold(n_splits=5, random_state=100)
model_skfold = LogisticRegression()
results_skfold = cross_val_score(model_skfold, X, Y, cv=skfold)
print("STRATIFIED K-FOLD SCORE: %.2f%%" % (results_skfold.mean()*100.0))
```

STRATIFIED K-FOLD SCORE: 64.50%
/usr/local/lib/python3.7/dist-packages/sklearn/model_selection/_split.py:296: FutureWarning: Setting a random_state h
FutureWarning

▼ Cumulative Accuracy Profile (CAP) Curve

```
[ ] total = len(y_test)
    class_1_count = np.sum(y_test)
    print(class_1_count)
    class_0_count = total - class_1_count
    plt.plot([0, total], [0, class_1_count], c = 'r', linestyle = '--', label = 'Random Model')
```

✓ 0s completed at 19:50

Day_9_Validating_Model_Ad_Sale x

colab.research.google.com/drive/1mRUJ74j67HCZszOODXsUz4Z-uVTNtmB-#scrollTo=DcsG6RJqacbw

Apps (2) SLAM for the ro... Sensors - ROS Wiki (2) [ROS Q&A] 031... The Secret of Nikol... (2) How to Set Up T... (8) Setting Up IR re... jetson.pdf How to recover the... fastai_deeplearn_pa... Reading list

Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_...

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text


```
plt.plot([0, class_1_count, total],
         [0, class_1_count, class_1_count],
         c = 'grey',
         linewidth = 2,
         label = 'Perfect Model')

probs = model.predict_proba(X_test)
probs = probs[:, 1]
model_y = [y for _, y in sorted(zip(probs, y_test), reverse = True)]
y_values = np.append([0], np.cumsum(model_y))
x_values = np.arange(0, total + 1)

plt.plot(x_values,
         y_values,
         c = 'b',
         label = 'LR Classifier',
         linewidth = 4)
```

RAM Disk Editing

0s completed at 19:50



Day_9_Validating_Model_Ad_Sale_Prediction_from_Existing_customer_... ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

RAM Disk Editing

```
index = int((50*total / 100))


## 50% Vertical line from x-axis
plt.plot([index, index], [0, y_values[index]], c='g', linestyle = '--')

## Horizontal line to y-axis from prediction model
plt.plot([0, index], [y_values[index], y_values[index]], c = 'g', linestyle = '--')

class_1_observed = y_values[index] * 100 / max(y_values)
plt.xlabel('Total observations')
plt.ylabel('Class 1 observations')
plt.title('Cumulative Accuracy Profile')
plt.legend(loc = 'lower right')
```

100
39
<matplotlib.legend.Legend at 0x7feb240feb10>
Cumulative Accuracy Profile

✓ 0s completed at 19:50



<matplotlib.legend.Legend at 0x7feb240feb10>

